

Current-Limiting Motor Control

Chris Gerth - Jeremy Lee - FRC 1736 Robot Casserole

Dec 2015

Abstract

A method for limiting motor voltage commands is presented. The target controlled variable is system voltage - too much current draw from the motors will lower system voltage to the point where the controller enters a “brownout” state. This is undesirable during a FRC robotics match. The method is an adaptive observer plant model which estimates battery parameters in real time, and limits current draw from motors to prevent system voltage drop.

Contents

I	Problem Statement, Constraints, and Assumptions	3
II	Current Limiting Algorithm	3
1	Data Sources	4
2	Motor Model	4
3	Determining Motor Constants	5
4	Total Estimated Current Draw	6
5	Battery Model	6
6	Limiting Method	7
III	Battery Parameter Estimation Algorithm	8
7	Data Sources	8
8	Estimation Method	9
8.1	Calculations	12
8.1.1	Averaging Filters	12
8.1.2	R_{bat} Calculation using Least Mean Squares on a Set of Points	12
8.1.3	Confidence Calculation and Filtering	13
8.1.4	V_{oc} Calculation	13
9	Usage of Estimated Parameters	13
IV	Implementation Details	13

Part I

Problem Statement, Constraints, and Assumptions

During FRC robotics competitions, there are many constraints on the set of electrical components which may be used. This includes motors, controllers, and batteries. The construction of an ever-more powerful drive-train is inherently limited by these component constraints, as arbitrarily more-powerful motors and energy sources cannot be used. The usual standard for drive-train combinations involves the FR801-001 2.5" CIM motor and the MK ES17-12 battery. There are a maximum of six CIM motors used to power all wheels. The motors are fed by electronic PWM controllers, where voltage is the commanded parameter. CIM motors can pull upwards of 120A at stall. Six stalled CIM motors is more than enough to cause controller brownout due to current-draw-induced voltage drop (verily, six stalled CIM motors can cause the MK battery to combust! This is why circuit breakers exist!). Even at non-stall conditions, allowing the motors to draw excessive current can quickly pull down system voltage. Traditionally, this is limited by using fewer CIM motors for the drive-train, or using less grippy wheels (allowing for slip to occur sooner), or increasing the gear ratio to slow the whole system down. All of these mitigations reduce drive-train performance, and are therefore considered to be an engineering trade-off.

The low-bar solution for limiting system voltage drop is to monitor either the system voltage or battery current draw, and reducing all motor commands as a preset limit is neared. This may work in some cases. However, most brownouts occur due to transient conditions (IE, current spikes) from collisions or sudden changes in driver commands. Since the measurements of system voltage and current are analog, they are inherently limited both by the system sample time, and by analog noise. Getting accurate signals will involve filtering, but any filtering induces delay in signal response time to rapidly changing conditions. A higher-bandwidth data source is needed. For this algorithm, the speed of the motors was chosen, since encoders have much more accurate signals exactly at the sample rate of the software (nearly zero noise). This leads to the need for an observer plant model to estimate motor current draw based on speed, voltage, and other physical parameters of the robot. Finally, the traditional model of ideal DC motors yields a total current draw, but does not directly predict system voltage drop due to this current draw. Predicting system voltage drop requires knowing an accurate model of the battery. Since each battery is slightly different and over a match the battery's parameters will change, an adaptive algorithm is needed to estimate the battery's parameters. Since these parameters change relatively slowly over time, lower-bandwidth signals may be used as inputs to derive them.

This paper presents a method which allows for optimal mechanical design of the drive-train, but actively ensures system voltage does not drop below a preset minimum threshold (referred to herein as $V_{sys,min}$). By limiting the system voltage drop in software, the drive-train can be designed to be arbitrarily powerful and grippy, and allow for as rapid of a gear ratio as desired. Since the algorithm only requires a preset $V_{sys,min}$ limit (along with a few well-known physical parameters of the motors), it is guaranteed to always run the hardware at its physical limits, and not impose additional "tuning" constraints seen by other similar algorithms.

For the purposes of this paper, a six-CIM tank drive platform is assumed. This means two sets of three-motor drive-trains with a single gear ratio from motor to wheel, which apply force on the left and right sides of a rigid frame. Rotation is accomplished by driving the wheels at different speeds. However, the algorithms described herein are easily adapted to other drive platforms.

Part II

Current Limiting Algorithm

The Current Limiting Algorithm is responsible for establishing a maximum motor voltage command to each set of wheels. The limiting involved is driven by the need to prevent system voltage from dropping below the preset $V_{sys,min}$ limit.




Figure 1: Classical model of a DC Motor

1 Data Sources

As indicated in Part I, the only external input needed for current estimation is the motor's present rotational speed. This may be derived from an encoder attached somewhere on each side of the drive-train. The motor's signed rotational speed at time-step n in rad/sec can be calculated as

$$\omega_m[n] = K_{enc_ratio} * \omega_{enc}[n] \quad (1)$$

Here, ω_{enc} is the encoder's rotational speed in radians/sec at time-step n , and K_{enc_ratio} is the gear ratio between the motor, and the measuring point of the encoder. If the encoder is attached directly to the motor, K_{enc_ratio} is simply 1. If the encoder is attached on another gear, it will be something other than 1.

Additionally, we will need to know the voltage applied to the motor. This can be calculated based off the present system voltage and the software command to the motor. The software command is normalized to the range $[-1, 1]$, where -1 means "full reverse" and 1 means "full forward" (0, of course, means "stop"). The present system voltage can come from a number of places: For simplicity it can be assumed to just be 12V. It could also be taken from a filtered measurement of the present system voltage (which will be needed anyway in part III). Finally, it could also be derived from the previous loop estimation from this current limiting algorithm. The first solution does not account for decaying battery voltage over the match, but the final one may induce oscillation due to an extra degree of linkage (higher order) in the differential equations. Therefore, for this analysis, the filtered measurement of present system voltage will be used.

$$V_m[n] = V_{sys}[n] * Cmd[n] \quad (2)$$

2 Motor Model

The classical model for a DC motor involves a two-terminal device circuit, where the two terminals are linked by a resistance and a variable voltages source in series. The resistance represents the electrical resistance of the motor winding and communicator. All inductive effects of the wrapped wire are ignored. The voltage source represents the "back-EMF" induced by the fact that a magnet is spinning with respect to a coil of wire. The current through the series circuit is proportional to the torque of the motor. The speed of the motor shaft is proportional to the back-EMF effect of the voltage source.

Assuming V is the driving voltage from the motor controller and L is zero, we can apply Kirchhoff's Voltage Law and Ohm's Law around the circuit to arrive at the following relationship:

$$V_m[n] = I_m[n]R_m + V_{emf}[n]$$

For V_m being the driving voltage of the motor, R_m being the resistance of the motor winding, I_m being the current draw of the motor, and V_{emf} being the back-EMF from the output shaft rotation.

Re-arranging, we find the current draw from one motor on one side of the drive-train is:

$$I_m[n] = \frac{V_m[n] - V_{emf}[n]}{R_m}$$

Again from the classical DC motor model, we assume that V_{emf} is linearly proportional to the rotational speed of the motor. We will call this constant of proportionality K_i . This yields the final relationship

$$I_m[n] = \frac{V_m[n] - K_i \omega_m[n]}{R_m} \quad (3)$$

We have now reduced the total current from one motor to an equation made up of known inputs and constant values.

3 Determining Motor Constants

For the CIM motor, we know two crucial facts: Stall Current is 133A, and Free-wheel speed is 5310 RPM while drawing 2.7A¹. In both cases, the supply voltage is 12V. Using these relationships, we can solve for the constant parameters in the motor model. Starting with Stall condition, we know $\omega_m = 0$. Therefor:

$$\begin{aligned} I_m[n] &= \frac{V_m[n] - K_i \omega_m[n]}{R_m} \\ 113 &= \frac{12 - K_i * 0}{R_m} \\ R_m &= \frac{12}{113} \\ R_m &= 0.1062\Omega \end{aligned} \quad (4)$$

Now, using this number, we plug in information for Free-Wheel Speed. Remember to convert RPM to rad/sec

$$\begin{aligned} I_m[n] &= \frac{V_m[n] - K_i \omega_m[n]}{R_m} \\ 2.7 &= \frac{12 - K_i * 5310 * \frac{\pi}{180}}{0.1062} \\ K_i &= 0.1263 \end{aligned} \quad (5)$$

Remember these constants are for CIM motors only. Different motors will need this section to be re-calculated for their parameters.

¹See <http://content.vexrobotics.com/docs/217-2000-CIM-motor-specs.pdf>




Figure 2: Classical model of a lead-acid Battery

4 Total Estimated Current Draw

Since all parameters of the motors are now known, we can use the speed from either side of the drive-train to determine total drive-train current draw. Again, note this assumes a 6-CIM tank drive setup, with subscript r and l indicating left and right sides of the drive-train:

$$I_{dt}[n] = 3 * I_{mr}[n] + 3 * I_{ml}[n]$$

$$I_{dt}[n] = 3 * \frac{V_{mr}[n] - K_i \omega_{mr}[n]}{R_m} + 3 * \frac{V_{ml}[n] - K_i \omega_{ml}[n]}{R_m} \quad (6)$$

Again, note this equation for total current draw produces the current at time-step n using only known inputs (V and ω) and constant parameters.

At this point, limiting could be applied based around a maximum desired current draw from the drive-train. However, since brownouts are caused by system voltage drops, a further step in this algorithm is taken to estimate the system voltage drop induced by this current draw.

5 Battery Model

Before continuing determining if limiting is actually needed, we must establish what our battery behaves like under load. A battery is classically characterized as an ideal voltage source in series with a resistor. As more current is drawn from the battery, the voltage drop across the resistance causes the output voltage of the battery to drop. As the battery discharges, the ideal voltage source decreases from a nominal (fully-charged) value of near $\sim 13V$ to something much lower ($\sim 8V$ or even lower). As the battery ages over time, the internal resistance will also increase, making the battery more susceptible to voltage drops from a large current draw. See figure 2.

From section 4, we know during every control loop what the current draw from the drive-train will be if we apply the driver-desired voltage to the motors at the drive-train. Given this current draw $I_{dt}[n]$, we can then estimate the system voltage drop via the following equation:

$$V_{sys}[n] = V_{oc} - R_{bat}I_{dt}[n] \quad (7)$$

For V_{sys} being the system voltage at time-step n , V_{oc} being the open-circuit voltage of the battery (the ideal voltage source in figure 2), R_{bat} being the battery's internal resistance, and I_{dt} being the current draw from the robot. This is assumed to be drive-train-only, all other current sinks on the robot are negligible. Note the battery's open-circuit voltage and internal resistance are considered constants for this part of the analysis. For reference, they are usually around $V_{oc} = 12$ and $R_{bat} = 0.012\Omega$ for a healthy battery.² However, later in the analysis, it will be shown how to calculate them deceptively over time.

6 Limiting Method

Limiting system voltage drop involves a multi-step process. The basic algorithm is:

1. Estimate current battery parameters
2. Measure motor speeds from each side of the drive-train
3. Determine the driver demanded voltages (V_{ddr}, V_{ddl}) to the drive-train
4. Estimate the current draw from the drive-train **if** the driver demanded voltages were sent to the motors on this time-step
5. Estimate the system voltage given the estimated current draw from the drive-train
6. **If** the estimated system voltage is below the preset threshold $V_{sys,min}$, calculate a scaling factor for the driver-demanded voltages which will hold V_{sys} at $V_{sys,min}$.
7. **Else**, set the driver-demanded voltage to the motor (no limiting)

Every step except 6 has been demonstrated already. The crucial task now is determining a scaling factor (which we will call γ) to apply to the driver-demanded voltages.

Estimating the system voltage drop can be done by combining equations 6 and 7:

$$V_{sys,est}[n] = V_{oc} - R_{bat} \left[3 * \frac{|V_{ddr}[n] - K_i\omega_{mr}[n]|}{R_m} + 3 * \frac{|V_{ddl}[n] - K_i\omega_{ml}[n]|}{R_m} \right] \quad (8)$$

For $V_{sys,est}$ being the estimated system voltage for driver-demanded motor voltages V_{ddr} and V_{ddl} . Note the introduction of absolute-value signs to account for the fact that the current direction is inverted at the motor controller, so the battery always sees current drawn out of it (even when going in reverse).

If scaling is needed (IE, $V_{sys,est} < V_{sys,min}$), we plug in the scaled values to equation 8 and solve for the scaling value γ . Starting from equation 8 and plugging in the known values for the "scaling needed" condition:

$$\begin{aligned} V_{sys,min} &= V_{oc} - R_{bat} \left[3 * \frac{|\gamma V_{ddr}[n] - K_i\omega_{mr}[n]|}{R_m} + 3 * \frac{|\gamma V_{ddl}[n] - K_i\omega_{ml}[n]|}{R_m} \right] \\ V_{oc} - V_{sys,min} &= R_{bat} \left[3 * \frac{|\gamma V_{ddr}[n] - K_i\omega_{mr}[n]|}{R_m} + 3 * \frac{|\gamma V_{ddl}[n] - K_i\omega_{ml}[n]|}{R_m} \right] \\ \frac{V_{oc} - V_{sys,min}}{3R_{bat}} &= \frac{|\gamma V_{ddr}[n] - K_i\omega_{mr}[n]|}{R_m} + \frac{|\gamma V_{ddl}[n] - K_i\omega_{ml}[n]|}{R_m} \\ \frac{R_m (V_{oc} - V_{sys,min})}{3R_{bat}} &= |\gamma V_{ddr}[n] - K_i\omega_{mr}[n]| + |\gamma V_{ddl}[n] - K_i\omega_{ml}[n]| \end{aligned}$$

²See <http://www.mkbattery.com/images/ES17-12.pdf>

Due to the double absolute value symbols, we now have four possible solutions for γ . Doing some funky algebra³, we find all four solutions for gamma:

$$\gamma = \begin{cases} -\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) - (K_i\omega_{ml}[n]) \\ \frac{V_{ddr}[n] - V_{ddl}[n]}{\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) - (K_i\omega_{ml}[n])} \\ -\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) + (K_i\omega_{ml}[n]) \\ \frac{V_{ddr}[n] - V_{ddl}[n]}{\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) + (K_i\omega_{ml}[n])} \\ -\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) - (K_i\omega_{ml}[n]) \\ \frac{V_{ddr}[n] + V_{ddl}[n]}{\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) - (K_i\omega_{ml}[n])} \\ -\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) + (K_i\omega_{ml}[n]) \\ \frac{V_{ddr}[n] + V_{ddl}[n]}{\left(\frac{R_m(V_{oc}-V_{sys,min})}{3R_{bat}}\right) + (K_i\omega_{mr}[n]) + (K_i\omega_{ml}[n])} \end{cases} \quad (9)$$

All four possible values for γ should be computed at run-time. From the four possible values, it is known that gamma can be in the range $[0, 1]$ since the scaling cannot exceed the physical limits of what was requested or is possible with the motor controllers. Among all the solutions which are in this range, the largest should be chosen (as the $\gamma = 1$ case is where the driver-demanded voltage is honored). If no solution is within the $[0, 1]$ range, default to $\gamma = 0$. This assumes current cannot be back-driven through an active controller, and the braking effect of the zero-speed-commanded controller will push the system back to steady-state without any external control effort.

Once a suitable γ is found, the motor voltages for each motor should be applied as such:

$$V_m = \begin{cases} V_{dd} & V_{sys,est} \geq V_{sys,min} \\ \gamma V_{dd} & V_{sys,est} < V_{sys,min} \end{cases} \quad (10)$$

Part III

Battery Parameter Estimation Algorithm

Up until now, it has been assumed that the battery parameters are constants ($V_{oc} = 12$ and $R_{bat} = 0.012\Omega$). This is valid for some batteries, but will induce significant error as batteries discharge over a match, and age over many seasons. While the aforementioned numbers are good starting points, it is worthwhile attempting to determine these parameters on the fly, to account for different conditions. Since they are slowly-changing compared to driver inputs and motor speeds, the previous calculations will be unaffected (as the assumption that the battery parameters are “constant” is still fairly true with respect to the other time-varying signals). The basic estimation method will use known voltage and current values over multiple time-steps to calculate parameters of the simplified battery model presented in figure 2.

7 Data Sources

The battery model has two degrees of freedom (Open-Circuit voltage and internal series resistance), so two sources of data are needed to fully constrain the system. Starting in 2015, a CAN-enabled Power Distribution Panel was made standard for FRC. Using the API's associated with it, it is possible to read the total current from the battery, and the present system voltage. As mentioned already, these signals are lower-bandwidth and would not be good for estimating transient values. However, outside of extreme transients, their values can be filtered down to fairly accurate values for the voltage and current just at the battery's connection to the robot.

The measured current and voltage values for the system form points on the I-V plot for the battery. It should be noted that the slope of this plot ends up being the equivalent series resistance (R_{bat}), and the y-intercept is the open-circuit battery voltage (V_{oc}). The basic method for determining these two desired parameters is to take many system current and voltage measurements, and use them to create a best-fit line. The slope and y-intercept of this best-fit line are then used to calculate the V_{oc} and R_{bat} parameters needed in the current-limiting algorithm. To get

³See www.wolframalpha.com

an accurate best-fit line, it is required that the current and voltage measurements be spread out over the I-V plane. In other words, a steady-state robot sitting still will not produce meaningful measurements - it is required to cycle the motors/compressors on and off to generate a changing current draw from the battery. R_{bat} on particular is sensitive to steady-state error due to noise. While the current is changing, the estimated R_{bat} parameter can be trusted, but a sample-and-hold methodology must be used when it is not changing much.

8 Estimation Method

Figure 2 is referenced as the model for the battery, with the measured system voltage and measured system current present at its output ports.

The estimation method can be summarized as follows:

- Read in the present system voltage and current
- Average some amount of previous data to help reject noise in the measurement.
 - Simulated experiments showed 300ms of data averaging was a good starting point. Make the length of the averaging tunable, and adjust to taste on the actual robot.
 - * At 10ms, this equates to the past 30 samples
 - This will generate two filtered signals - one for the system current, one for the system voltage
- Consider a window of the filtered data stretching from the present into the past for a given length of time
 - Simulated experiments showed 1000ms for a window length worked well
- Estimate the equivalent series resistance of the battery from this window of samples
 - See section 8.1.2 for the algorithm
- Calculate the spread of current draw readings in the window under consideration
 - Here, we define spread as the standard deviation of the points within the window
- If the spread is above a constant, tunable threshold, use the estimated ESR in all other calculations, assuming it to be correct.
 - Additionally, if this is the biggest spread seen since the last time spread went above the threshold, record the spread and ESR values for later use.
- If the Spread is below the threshold, use the previous best-spread ESR. (sample and hold most-confident value)
 - Reset the “largest spread” threshold to zero as well in prep for the next time spread goes above the threshold

The results of such an algorithm can be seen in the provided figures: Note that when confidence goes to “1”, the percent error goes to nearly zero and the estimated ESR “catches up” with the actual one. When information is sparse in-between large current draw changes, the best last known-good value for ESR is held. This is a reasonable approximation, as percent error stays mostly below 10% for both ESR and Voc. The figures illustrate the algorithm working in the presence of noise in the readings for system voltage and current.




Figure 3: Performance of Algorithm on a pulsed-current-draw waveform with a discharging battery

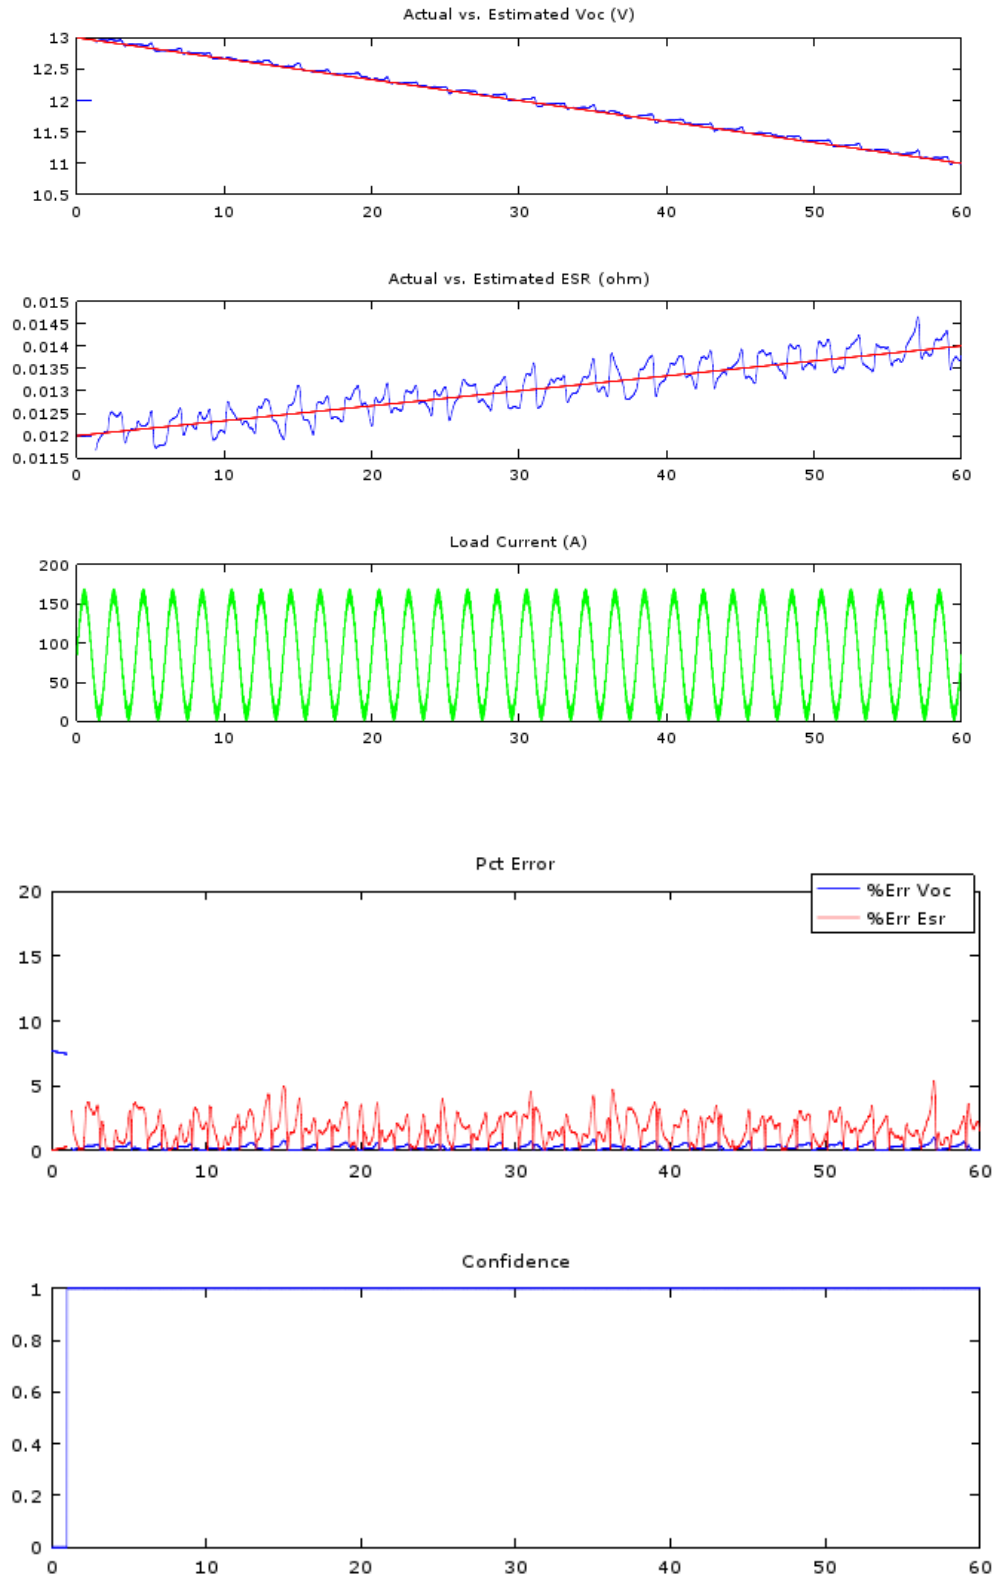


Figure 4: Performance of Algorithm on a sine waveform current draw with a discharging battery

8.1 Calculations

8.1.1 Averaging Filters

The algorithm refers to utilizing an averaging filter to eliminate some of the noise from the measured input I_{sys} and V_{sys} readings. An averaging filter has one input and one output - the output is always equal to the average (or arithmetic mean) of the last N inputs, where N is said to be the “length” of the filter. For example, on the system voltage,

$$V_{sys}[n] = \frac{\sum_{i=0}^{N-1} V_{meas}[n-i]}{N}$$

Where V_{sys} ends up being the filtered voltage value at time n , and V_{meas} is the voltage value read from the PDB sensor via a function call in software. N 's value can be determined by the length of averaging needed and the sample time. If it is desired to average 2 seconds of data, and the sample time is 10 ms, you will need 200 points ($N = 2/Ts = 200$).

Use averaging filters on both the read-in system current and voltage. It may also be used on the ESR calculated.

8.1.2 R_{bat} Calculation using Least Mean Squares on a Set of Points

From <http://faculty.cs.niu.edu/~hutchins/csci230/best-fit.htm>.

The algorithm used to calculate R_{bat} from a window of current and voltage measurement points is based on the “Least-Mean-Squares” algorithm. LMS seeks to define a line with a slope such that the distance from the line to any point in the set is as small as possible. That is to say, the line is the closest to the average of all the points as much as possible. The algorithm is derived by first defining this “best fit” condition mathematically, and then working backward to the line equation. That proof is not covered in this paper. Instead, the final algorithm is presented.

Define first N_{lms} to be the number of points in the window of I-V readings to consider. This is referred to as the “window size”, and is numerically the same as the number of current-voltage pairs plotted on the I-V graph from which we will attempt to apply a best-fit line.

Define a few more variables:

$$\begin{aligned} S_v[n] &= \sum_{i=0}^{N_{lms}-1} V_{sys}[n-i] \\ S_I[n] &= \sum_{i=0}^{N_{lms}-1} I_{sys}[n-i] \\ S_{IV}[n] &= \sum_{i=0}^{N_{lms}-1} (I_{sys}[n-i] * V_{sys}[n-i]) \\ S_{I^2}[n] &= \sum_{i=0}^{N_{lms}-1} (I_{sys}[n-i]^2) \\ \bar{V}[n] &= \frac{S_v}{N_{lms}} \\ \bar{I}[n] &= \frac{S_I}{N_{lms}} \end{aligned}$$

Then, using these variables, the ESR estimate for timestep n may be calculated:

$$R_{bat}[n] = -\frac{S_{IV}[n] - (S_I[n] * \bar{V}[n])}{S_{I^2}[n] - (S_I[n] * \bar{I}[n])}$$

8.1.3 Confidence Calculation and Filtering

As stated, the confidence we have in the estimated $R_{bat}[n]$ is based off of the standard deviation of the set of current measurements within the window being considered. For reference, this standard deviation (or “spread”) is:

$$\sigma_{conf}[n] = \sqrt{\frac{1}{N_{lms}} \sum_{i=0}^{N_{lms}-1} \left((I_{sys}[n-i] - \bar{I}[n])^2 \right)}$$

A tunable constant σ_{min} determines the minimum spread needed before the R_{bat} estimation is trusted. If $\sigma_{conf}[n]$ is larger than the minimum, use $R_{bat}[n]$. Otherwise, re-assign $R_{bat}[n]$ to the value during the highest- $\sigma_{conf}[n]$ loop during the most recent period of trustable- R_{bat} numbers. A snippet of .m code is provided illustrating this logic:

```
% If the spread is above a tuned minimum threshold, we may use this window
% for the ESR calculation.
if(spread > min_spread_thresh_A)
    confidence(i) = 1;
    % Additionally, if this is the largest spread we've seen so far,
    % save the ESR value for when the spread is no longer
    % big enough to trust the calculation.
    if(spread > prev_best_spread)
        prev_best_spread = spread;
        prev_best_esr = ESR_est_raw(i);
    end

else
    %When the spread isn't big enough, we don't trust our calculation
    % We also reset the "best spread" to zero since we are no longer in a
    % "confidence = 1" region.
    confidence(i) = 0;
    prev_best_spread = 0;
end

%If we didn't trust this window's calculation, use the previous best calculation
if(confidence(i) == 0)
    ESR_est_raw(i) = prev_best_esr;
end
```

8.1.4 V_{oc} Calculation

Regardless of whether the R_{bat} calculation was trusted or not, the open-circuit voltage can always be calculated via this formula:

$$V_{oc}[n] = \bar{V}[n] + R_{bat}[n] * \bar{I}[n]$$

9 Usage of Estimated Parameters

The two battery parameters have just been estimated from measured system voltage and current values. This calculation should be done as part of the first step of the limiting algorithm. The determined numbers for the open-circuit voltage and internal resistance of the battery may then be used in the other calculations.

Part IV

Implementation Details

The equations presented above are in a verbose form to demonstrate their physical meaning. Simplifying them is very possible to reduce the number of computations done at run-time.

Additional voltage drops and resistances may need to be added to model the losses in the wires and motor controllers.

It is highly recommended to keep a log of pertinent internal information. Included in this is the measured system voltage and current, estimated system voltage and drive-train currents, the estimated battery parameters, the utilized scaling factor γ , as well as the driver-demanded and commanded voltages. These logs can be used to validate proper software behavior, and diagnose erratic motions. Additionally, any times when limiting was applied can be used to show why exactly the algorithm was needed.

Finally, it is to be noted that this algorithm will necessarily reduce the apparent responsiveness of the drive-train. Hopefully the 6-CIM drive is powerful enough that the reduction from this algorithm will not be missed. However, it is ideal to have it in place well before the drive team begins practice, as adding it partway in may be perceived as an undesired limitation.